



Speech Recognition





Speech Recognition

- Fundamentals of Digital Speech Processing
- Mel-Frequency Cepstral Coefficients (MFCCs)
- Speech Recognition by Dynamic Time Warping
- Speech Recognition by Hidden Markov Models

Human Speech Production

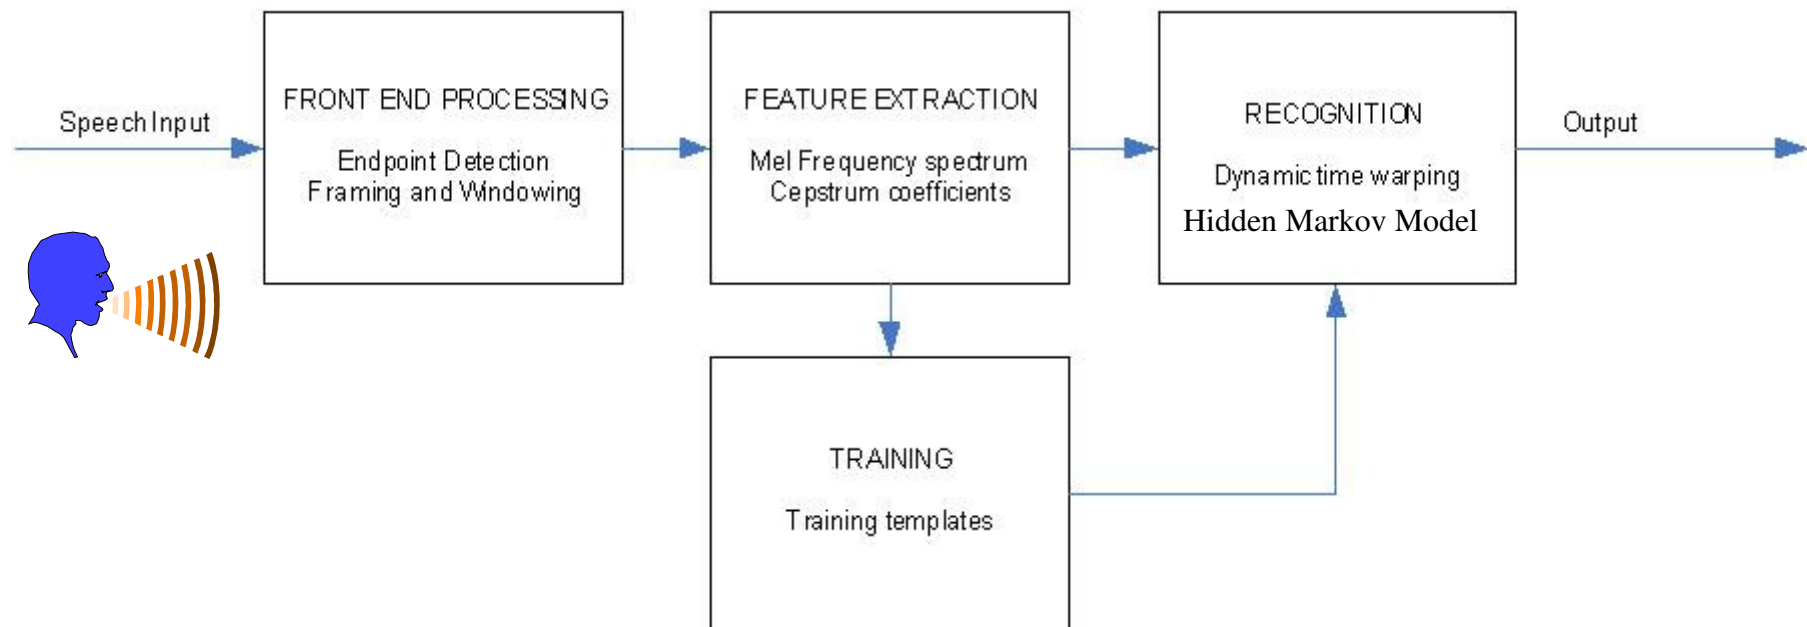
- All speech sounds are formed by blowing air from the lungs through the vocal tract.



Automatic Speech Recognition (ASR) Paradigms

- Continuous vs. Isolated
- Large (>1000) vs. Small Vocabularies (< 100)
- Speaker Dependent vs. Speaker Independent
- Speech Recognition vs. Speaker Recognition
- Speaker Recognition vs. Speaker Verification
- Context Dependent vs. Context Independent Verification
- Key Word Spotting
- SubWord Speech Units and Modeling
- Statistical Language Modeling and Perplexity
- Robust Speech Recognition and Adaptation

A Typical ASR System




Why ASR is Difficult?

- Speech signals are **continuous**.
 - No explicit markers to indicate end of one sound and start of the next.
- Speech signals are **highly variable**.
 - Not only differences as a result of different people/sex saying the same word/sentence, but also differences with the **same** person saying the **same** word/sentence at different times.
- Speech is **ambiguous**.
 - There is no acoustic difference between **to**, **two** and **too**.
- Speech is **contaminated**.
 - -- Usually a speech signal occurs in an environment where there is some degree of reverberation, or competing acoustic noises

MFCC (see speaker_recognition project)



- Mel-Frequency Cepstral Coefficient (MFCC)
 - Most widely used spectral representation in ASR



Pre-Emphasis before DFT

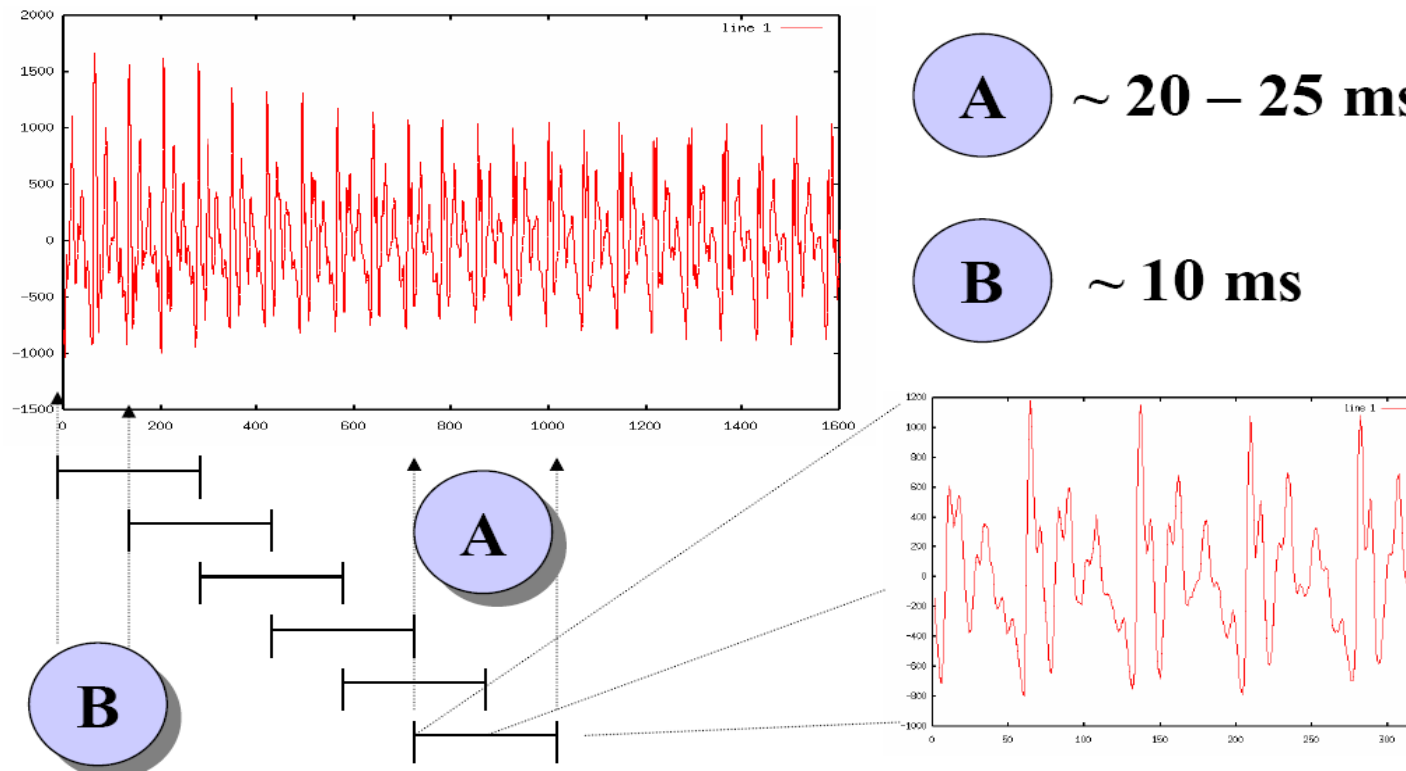
- **Pre-emphasis:** remove the lowpass lip radiation effect and boost the energy in the high frequencies
- Boosting high-frequency energy gives more info to Acoustic Model – better recognition

$$1 - \alpha z^{-1}, \quad \alpha = 0.97$$




Windowing (Framing)

- Apply **Hamming window**, duration 200 samples (25 msec) every 10 ms (100-Hz frame rate)




Quasi-Stationary Signal

Hamming Windowed Frames



(a) Rectangular window




(c) Hamming window

Mel-Scale & Mel Filter Bank

- Human hearing is not equally sensitive to all frequency bands, **less sensitive at higher frequencies**
- Human perception of frequency is **non-linear**
- Mel-scale is approximately linear below 1 KHz and logarithmic above 1 KHz

$$Mel(f) = 2595 \log_{10} \left(1 + \frac{f}{700} \right)$$



Mel-filter Bank Processing


- (20-24) triangular shaped filters spaced evenly along the Mel Frequency Scale with 50% overlap
- Energy from each filter is computed (N = DFT size, P = #filters) at time t:

$$e[j][t] = \sum_{k=0}^{N-1} H_j[k] \cdot \left\| \tilde{S}_t[k] \right\|^2 \quad \text{for } j = 1 \dots P$$


Triangular Filter *Signal Power Spectrum*

Equally Spaced Filters

- **Equally spaced** filters along the Mel-frequency scale with 50% overlap



- Analogous to **non-uniformly** spaced filters along linear frequency scale:



DCT to Approximate IDFT

- Compute Log-Energies from each of P filters
- Apply Discrete Cosine Transform (DCT)

$$MFCC[i][t] = \sqrt{\frac{2}{P}} \sum_{j=1}^P \left\{ (\log e[j][t]) \cdot \cos\left(\frac{\pi i}{P} (j - 0.5)\right) \right\}$$

- DCT: (1) improves diagonal covariance assumption, (2) compresses features
- Typically 12-14 MFCC features are extracted (higher order MFCCs useful for speaker-ID)

Distance Measure in the Pattern Matching

- Given one frame of speech (testing):

$$\vec{O} = (O_1, O_2, O_3, \dots, O_{12}),$$

- and another frame of speech (template):

$$\vec{E} = (E_1, E_2, E_3, \dots, E_{12}).$$

- the distance between these two frames of speech is (can be weighted if not cepstral coeff.):

$$d = dist(\vec{O}, \vec{E}) = \sum_{m=1}^{12} (O_m - E_m)^2$$

- Note that time domain comparison is meaningless, especially human ears are insensitive to time delay and slight vocal tract variations, which result in big changes in time domain.

Endpoint Detection

- To determine the **beginning** and the **end** of an isolated utterance (isolated word, isolated sentence, etc).
- Two factors: **energy** (sum of the magnitude in each frame) and the zero crossing rate (**ZCR**).
- For either “energy” or “ZCR”, there is an associated “**possible threshold**” P_T , and “**word start threshold**” WS_T .
- The P_T is “set” when the energy is just above background noise level (may be exceeded by spurious noise). The WS_T is “set” when the system is sure a word is spoken.
- Two additional thresholds: **minimum word length threshold**, ML_T (i.e., minimum number of frames per word), and **minimum silence duration threshold** MD_T , (i.e., minimum number of frames after the end of a word).

Start of A Word

- The condition: $frame - energy > WS_T^{(E)}$, or $ZCR > WS_T^{(Z)}$
- Once $P_T^{(E)}$ or $P_T^{(Z)}$ is exceeded, we then have to continue search forward until the corresponding WS_T is exceeded. If before WS_T is exceeded, P_T fails itself, then we have to restart the search.
- Once a valid word is detected, we start the search of the long-enough silence end MD_T . Once the end is identified, the word length must exceed the ML_T to be qualified as a word.

Pattern Matching

1. Dynamic Time Warping (DTW)
2. Hidden Markov Modelling (HMM)
3. Others



Template
MFCC

1	4	3	1	2	4	
3	2	1	5	7	5	
3	1	3	5	5	6	
3	5	8	2	4	5	
1	4	3	6	7	8	
2	1	7	5	8	6	



Test MFCC


DTW
Basics

It remains to find the warping function which minimizes

$$\sum_k d(c_k)w_k$$

A direct search is too slow.

We use the method known as Dynamic Programming.



Best path to (i,j) is the minimum of

- (i) best path to (i-1,j) +dist from (i-1,j) to (i,j)
- (ii) best path to (i-1,j-1) +dist from (i-1,j-1) to (i,j)
- (iii) best path to (i,j-1) +dist from (i,j-1) to (i,j)

Sakoe and Chiba algorithm

Dynamic programming solution to speech pattern matching problem:


$$f(1,1) = 2d(1,1) \quad (\text{assuming } c_0 = (0,0) \text{ so } w_1 = 2)$$

$$f(t,1) = f(t-1,1) + d(t,1) \quad (\text{for } 2 \leq t \leq T)$$

$$f(1,s) = f(1,s-1) + d(1,s) \quad (\text{for } 2 \leq s \leq S)$$

$$f(t,s) = \min \begin{cases} f(t-1,s) + d(t,s) & (\text{right}) \\ f(t-1,s-1) + 2d(t,s) & (\text{diagonal}) \\ f(t,s-1) + d(t,s) & (\text{up}) \end{cases} \quad (\text{for } 2 \leq t \leq T, 2 \leq s \leq S)$$

$$D(\underline{\mathbf{O}}, \underline{\mathbf{E}}) = f(T, S)$$



$f(t,s)$ is the shortest distance to the point (t,s)

$d(t,s)$ is distance between feature vectors \mathbf{O}_t and \mathbf{E}_s

Example

The bold numbers represent distances between feature vectors for input data and template, i.e. $d(t,s)$, the numbers in parentheses are the cumulative distances $f(t,s)$, and the arrows indicate the best path.

s	7 (16)	→	8 (24)		6 (24)		4 (26)
	↑				↑	↗	
	4 (9)	→	9 (18)		5 (18)	→	7 (25)
	↑				↑		
	3 (5)	→	7 (12)		3 (13)	→	7 (20)
	↑			↗			
	1 (2)	→	5 (7)	→	5 (12)	→	9 (21)
	t						

Minimum overall distance is 26. The optimum path can be found by tracing the arrows back from the top right corner.


Markov Models

- Set of **states**: $\{s_1, s_2, \dots, s_N\}$
- Process moves from one state to another generating a sequence of **states (observations)**: $s_{i1}, s_{i2}, \dots, s_{ik}, \dots$
- **Markov chain property**: probability of each subsequent state depends only on what was the previous state:

$$P(s_{ik} \mid s_{i1}, s_{i2}, \dots, s_{ik-1}) = P(s_{ik} \mid s_{ik-1})$$

- To define Markov model, the following probabilities have to be specified: transition probabilities $a_{ij} = P(s_i \mid s_j)$ and initial probabilities $\pi_i = P(s_i)$

Example of Markov Model



- **Two states** : 'Rain' and 'Dry'.
- **Transition probabilities**: $P(\text{'Rain'}|\text{'Rain'})=0.3$,
 $P(\text{'Dry'}|\text{'Rain'})=0.7$, $P(\text{'Rain'}|\text{'Dry'})=0.2$, $P(\text{'Dry'}|\text{'Dry'})=0.8$
- **Initial probabilities**: say $P(\text{'Rain'})=0.4$, $P(\text{'Dry'})=0.6$.

Calculation of Sequence Probability

- By Markov chain property, **probability of state sequence** can be found by the formula:

$$\begin{aligned}
 P(s_{i1}, s_{i2}, \dots, s_{ik}) &= P(s_{ik} \mid s_{i1}, s_{i2}, \dots, s_{ik-1}) P(s_{i1}, s_{i2}, \dots, s_{ik-1}) \\
 &= P(s_{ik} \mid s_{ik-1}) P(s_{i1}, s_{i2}, \dots, s_{ik-1}) = \dots \\
 &= P(s_{ik} \mid s_{ik-1}) P(s_{ik-1} \mid s_{ik-2}) \dots P(s_{i2} \mid s_{i1}) P(s_{i1})
 \end{aligned}$$

- Suppose we want to calculate a probability of a sequence of states (observations) in our example,

$$\begin{aligned}
 &\{\text{'Dry'}, \text{'Dry'}, \text{'Rain'}, \text{'Rain'}\}. \\
 &P(\{\text{'Dry'}, \text{'Dry'}, \text{'Rain'}, \text{'Rain'}\}) = \\
 &\quad P(\text{'Dry'}) P(\text{'Dry'} \mid \text{'Dry'}) P(\text{'Rain'} \mid \text{'Dry'}) \\
 &\quad P(\text{'Rain'} \mid \text{'Rain'}) = 0.6 * 0.8 * 0.2 * 0.3
 \end{aligned}$$

Hidden Markov Model (HMM)

- Set of states: $\{s_1, s_2, \dots, s_N\}$
- Process moves from one state to another generating a sequence of states :

$$s_{i1}, s_{i2}, \dots, s_{ik}, \dots$$

- Markov chain property: probability of each subsequent state depends only on what was the previous state:


$$P(s_{ik} | s_{i1}, s_{i2}, \dots, s_{ik-1}) = P(s_{ik} | s_{ik-1})$$

- States are not visible, but each state randomly generates one of M observations (or visible states)

$$\{v_1, v_2, \dots, v_M\}$$

- To define hidden Markov model, the following probabilities have to be specified: matrix of **transition probabilities** $A=(a_{ij})$, $a_{ij}= P(s_i | s_j)$, matrix of **observation probabilities** $B=(b_i(v_m))$, $b_i(v_m) = P(v_m | s_i)$ and a vector of **initial probabilities** $\pi=(\pi_i)$, $\pi_i = P(s_i)$. Model is represented by $M=(A, B, \pi)$.

Example of An HMM



Example of An HMM

- **Two states** : 'Low' and 'High' atmospheric pressure.
- **Two observations** : 'Rain' and 'Dry'.
- **Transition probabilities**:
 $P(\text{'Low'}|\text{'Low'})=0.3$, $P(\text{'High'}|\text{'Low'})=0.7$,
 $P(\text{'Low'}|\text{'High'})=0.2$, $P(\text{'High'}|\text{'High'})=0.8$
- **Observation probabilities** :
 $P(\text{'Rain'}|\text{'Low'})=0.6$, $P(\text{'Dry'}|\text{'Low'})=0.4$,
 $P(\text{'Rain'}|\text{'High'})=0.4$, $P(\text{'Dry'}|\text{'High'})=0.3$.
- **Initial probabilities**: $P(\text{'Low'})=0.4$, $P(\text{'High'})=0.6$.

Calculation of Observation Sequence Probability

- Suppose we want to calculate a probability of a **sequence** of observations in our example, **{‘Dry’, ‘Rain’}**.


- Consider all possible hidden state sequences:
$$P(\{\text{‘Dry’}, \text{‘Rain’}\}) = P(\{\text{‘Dry’}, \text{‘Rain’}\}, \{\text{‘Low’}, \text{‘Low’}\}) + P(\{\text{‘Dry’}, \text{‘Rain’}\}, \{\text{‘Low’}, \text{‘High’}\}) + P(\{\text{‘Dry’}, \text{‘Rain’}\}, \{\text{‘High’}, \text{‘Low’}\}) + P(\{\text{‘Dry’}, \text{‘Rain’}\}, \{\text{‘High’}, \text{‘High’}\})$$

where first term is :


$$\begin{aligned} &P(\{\text{‘Dry’}, \text{‘Rain’}\}, \{\text{‘Low’}, \text{‘Low’}\}) = \\ &P(\{\text{‘Dry’}, \text{‘Rain’}\} \mid \{\text{‘Low’}, \text{‘Low’}\}) P(\{\text{‘Low’}, \text{‘Low’}\}) = \\ &P(\text{‘Low’}) P(\text{‘Dry’} \mid \text{‘Low’}) P(\text{‘Low’} \mid \text{‘Low’}) P(\text{‘Rain’} \mid \text{‘Low’}) \\ &= 0.4 * 0.4 * 0.3 * 0.6 \end{aligned}$$

Speech Generation by HMM

Markov Generation Model



A Left-Right HMM Example



Three states + one starting state q^s + one final state q^f
 q^s and q^f are non-emitting states (**silence** is modeled here).

Assume there are 2 symbols to observe $X = \{x^1=a, x^2=b\}$

$$\Pi = \begin{bmatrix} 0.2 \\ 0.7 \\ 0 \\ 0.1 \end{bmatrix}$$

Initial state probabilities

$$A = \begin{bmatrix} 0 & 0.8 & 0.1 & 0.1 \\ 0 & 0 & 0.6 & 0.4 \\ 0 & 0 & 0.1 & 0.9 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Transition state probabilities

$$B = \begin{bmatrix} 0.8 & 0.2 \\ 0.4 & 0.6 \\ 0.1 & 0.9 \end{bmatrix} \quad \begin{matrix} P(a|q^1) \\ P(b|q^3) \end{matrix}$$


Observation symbol probabilities

Likelihood of An Observation Sequence

Given $X = \text{"aaa"}$, compute the likelihood for this model :
 $P(\text{aaa} \mid \lambda)$

The likelihood $P(X \mid \lambda)$ is given by the sum over all possible ways to generate X .

State sequence	Init	Obs a	Trans	Obs a	Trans	Obs a	Trans	Joint probability
$q^1 q^2 q^3$	0.2	0.8	0.8	0.4	0.6	0.1	0.9	0.0027648
$q^1 q^3 q^3$	0.2	0.8	0.1	0.1	0.1	0.1	0.9	0.0000144
$q^2 q^3 q^3$	0.7	0.4	0.6	0.1	0.1	0.1	0.9	0.0001512




$$P(\text{aaa} \mid \lambda) =$$

$$0.0029304$$

HMMs for Speech (Pattern) Recognition

Using HMM for speech recognition consists in computing the model λ_i among a set of K models which maximizes the likelihood for an observation to have been generated by this model :

$$\lambda_{\max} = \arg \max_{\lambda_i} P(X|\lambda_i) \quad \text{for } i = 1, \dots, K$$



Three Problems for HMMs

- Problem 1 : Recognition

Given $X = (x_1, x_2, \dots, x_T)$ and the various models $\{l_i\}$

How to efficiently compute $P(X|l_i)$?

Forward-Backward algorithm

- Problem 2 : Analysis

Given $X = (x_1, x_2, \dots, x_T)$ and a model l , find the optimal state sequence G . How can we undiscovered the sequence of states corresponding to a given observation ?

Viterbi algorithm

- Problem 3 : Learning

Given $X = (x_1, x_2, \dots, x_T)$, estimate model parameters $l = (P, A, B)$ that maximize $P(X|l)$, i.e., How do we determine the model parameters $l = (P, A, B)$?

Baum-Welch algorithm

References

- L. Rabiner, B.H. Juang, **Fundamentals of Speech Recognition**, Prentice Hall, 1993.
- L. Rabiner, B.H. Juang, “An introduction to hidden Markov models,” IEEE Signal Processing Magazine, 3(1):4-16, Jan 1986.